

# PARALLEL DISTRIBUTING COMPUTING



\*\_\_\_\_\_786  
\_\_\_\_\_\*

**Name: Muhammad Sharjeel Akhtar**

**Roll No: 20p-0101**

**Assignment No: 02**

**Submitted To Respected  
Sir: Dr Omer Usman Khan**

**Section: BCS-7F**

## Task 1: Getting Ready

P.S: Getting Login Issue With My Personal Account Which Was Working Fine Few Days Back So Using Friend Account For Testing Of Codes For The Assignment

Error:

```
spoofy@spoofy-Precision-M4600:~$ ssh pdc-p200101@121.52.146.108
Password:
Your account has expired; please contact your system administrator

Connection closed by 121.52.146.108 port 22
spoofy@spoofy-Precision-M4600:~$
```

## Doing SSH on Another Account:

```
spoofy@spoofy-Precision-M4600:~$ ssh pdc-p200045@121.52.146.108
Password:
pdc-p200045@lmar ~ $
```

SUCCESSFULLY LOGGED IN

## Task 2: Hello World

Initially create file locally in my computer,

```
spoofy@spoofy-Precision-M4600:~/Videos$ mkdir hello
spoofy@spoofy-Precision-M4600:~/Videos$ cd hello/
spoofy@spoofy-Precision-M4600:~/Videos/hello$ l
spoofy@spoofy-Precision-M4600:~/Videos/hello$ nano hello.cu
spoofy@spoofy-Precision-M4600:~/Videos/hello$ ls
hello.cu
```

Sending it to cloud pc,

```
^Cspoofy@spoofy-Precision-M4600:~/Videos/hello$ scp hello.cu pdc-p200045@121.52.146.108:/home/pdc-p200045/cuda/20p-0101
Password:
hello.cu
spoofy@spoofy-Precision-M4600:~/Videos/hello$
```

## VERIFICATION ON CLOUD:

```
pdc-p200045@lmar ~/cuda/20p-0101 $ ls
hello.cu
```

## RUNNING:

```
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc hello.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ls
a.out hello.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
hello world!
pdc-p200045@lmar ~/cuda/20p-0101 $ cat hello.cu
#include <stdio.h>
int main() {
    printf("hello world!\n");
    return 0;
}
pdc-p200045@lmar ~/cuda/20p-0101 $
```

## Task 5: Playing with 1D GPU indices

### Dummy Code that WE ARE GOING to USE:

```
/* Name: task5.cu
*/
#include <stdio.h>
__global__ void myHelloOnGPU(int *array) {
    // Position 1: To write Code here later
}
int main() {
    int N = 16;
    int *cpuArray = (int*)malloc(sizeof(int)*N);
    int *gpuArray;
    cudaMalloc((void **)&gpuArray, sizeof(int)*N);
    // Position 2: To write Code here later
    cudaMemcpy(cpuArray, gpuArray, sizeof(int)*N,
        cudaMemcpyDeviceToHost);
    int i;
    for (i = 0; i < N; i++) {
```

```
printf("%d ", cpuArray[i]);
}
printf("\n");
return 0;
}
```

## 1. Task 5a

```
Position 2: myHelloOnGPU<<<N, 1>>>(gpuArray);
Position 1: array[blockIdx.x] = blockIdx.x
```

### Result:

```
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task5a.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15
pdc-p200045@lmar ~/cuda/20p-0101 $
```

## 2. Task 5b

```
Position 2: myHelloOnGPU<<<N, 1>>>(gpuArray);
Position 1: array[blockIdx.x] = threadIdx.x
```

### Result:

```
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task5b.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
pdc-p200045@lmar ~/cuda/20p-0101 $
```

## 3.Task 5c

```
Position 2: myHelloOnGPU<<<1, N>>>(gpuArray);
Position 1: array[threadIdx.x] = threadIdx.x
```

**Result:**

```
pdg-p200045@lmar ~/cuda/20p-0101 $ ./a.out
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15
pdg-p200045@lmar ~/cuda/20p-0101 $
```

#### 4. Task 5d

```
Position 2: myHelloOnGPU<<<1, N>>>(gpuArray);
Position 1: array[threadIdx.x] = blockIdx.x
```

**Result:**

```
pdg-p200045@lmar ~/cuda/20p-0101 $ nvcc task5d.cu
pdg-p200045@lmar ~/cuda/20p-0101 $ ./a.out
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
pdg-p200045@lmar ~/cuda/20p-0101 $
```

#### 5. Task 5e

```
Position 2: myHelloOnGPU<<<1, N/2>>>(gpuArray);
Position 1: array[threadIdx.x] = threadIdx.x
```

**Result:**

```
pdg-p200045@lmar ~/cuda/20p-0101 $ nvcc task5e.cu
pdg-p200045@lmar ~/cuda/20p-0101 $ ./a.out
0 1 2 3 4 5 6 7 0 0 0 0 0 0 0 0
pdg-p200045@lmar ~/cuda/20p-0101 $
```

#### 6. Task 5f

```
Position 2: myHelloOnGPU<<<1, N/2>>>(gpuArray);
Position 1: array[threadIdx.x + blockDim.x] = threadIdx.x
```

**Result:**

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task5f.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
0 0 0 0 0 0 0 0 0 1 2 3 4 5 6 7
pdc-p200045@lmar ~/cuda/20p-0101 $

```

## 7. Task 5g

```

Position 2: myHelloOnGPU<<<N/2, 1>>>(gpuArray);
Position 1: array[blockIdx.x + gridDim.x] = 111 *(blockIdx.x + 1);

```

### Result:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task5g.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
0 0 0 0 0 0 0 0 0 111 222 333 444 555 666 777 888
pdc-p200045@lmar ~/cuda/20p-0101 $

```

## 8. Task 5h: What should be Position 1 and 2 in order to obtain the following output:

```

Position 2: myHelloOnGPU<<<N/2, 1>>>(gpuArray);
Position 1: array[blockIdx.x * 2] = 111 *(blockIdx.x + 1);

```

### Result:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task5h.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
111 0 222 0 333 0 444 0 555 0 666 0 777 0 888 0
pdc-p200045@lmar ~/cuda/20p-0101 $

```

## 9. Task 5j

```

Position 2: myHelloOnGPU<<<N, 1>>>(gpuArray);
Position 1: array[blockIdx.x] = gridDim.x - blockIdx.x - 1;

```

## Result:

```
pdcp200045@lmar ~/cuda/20p-0101 $ nvcc task5j.cu
pdcp200045@lmar ~/cuda/20p-0101 $ ./a.out
15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0
pdcp200045@lmar ~/cuda/20p-0101 $
```

## 10. Task 5k

```
Position 2: myHelloOnGPU<<N/4, N/4>>>(gpuArray);
Position 1: array[blockIdx.x * blockDim.x] = 111*(blockIdx.x + 1);
```

## Result:

```
pdcp200045@lmar ~/cuda/20p-0101 $ nvcc task5k.cu
pdcp200045@lmar ~/cuda/20p-0101 $ ./a.out
111 0 0 0 222 0 0 0 333 0 0 0 444 0 0 0
pdcp200045@lmar ~/cuda/20p-0101 $
```

## 11. Task 5m

```
Position 2: myHelloOnGPU<<N/4, N/4>>>(gpuArray);
Position 1: array[blockIdx.x * blockDim.x + threadIdx.x] =
111*(blockIdx.x + 1);
```

## Result

```
pdcp200045@lmar ~/cuda/20p-0101 $ nano task5m.cu
pdcp200045@lmar ~/cuda/20p-0101 $ nvcc task5m.cu
pdcp200045@lmar ~/cuda/20p-0101 $ ./a.out
111 111 111 111 222 222 222 222 333 333 333 333 444 444 444 444
pdcp200045@lmar ~/cuda/20p-0101 $
```

**Task 5n: What should be Position 1 and 2 in order to obtain the following output:**

```
Position 2: myHelloOnGPU<<<N/4, N/4>>>(gpuArray);  
Position 1: array[blockIdx.x * blockDim.x + threadIdx.x] = block  
threadIdx.x - 1;
```

**Result:**

```
pd-c-p200045@lmar ~/cuda/20p-0101 $ nvcc task5n.cu  
pd-c-p200045@lmar ~/cuda/20p-0101 $ ./a.out  
3 2 1 0 3 2 1 0 3 2 1 0 3 2 1 0  
pd-c-p200045@lmar ~/cuda/20p-0101 $
```

**Task 6: Playing with 2D GPU indices.**

**1. Task 6a:**

```
pd-c-p200045@lmar ~/cuda/20p-0101 $ nano task6a.cu  
pd-c-p200045@lmar ~/cuda/20p-0101 $ nvcc task6a.cu  
pd-c-p200045@lmar ~/cuda/20p-0101 $ ./a.out  
00 01 02 03  
04 05 06 07  
08 09 10 11  
12 13 14 15  
  
pd-c-p200045@lmar ~/cuda/20p-0101 $
```

**2. Task 6b:**



```

pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6a.cu
^[[Apdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
00 01 02 03
04 05 06 07
08 09 10 11
12 13 14 15

pdc-p200045@lmar ~/cuda/20p-0101 $ █

```

### 3. Task 6c:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6c.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
00 01 02 03
04 05 06 07
08 09 10 11
12 13 14 15

pdc-p200045@lmar ~/cuda/20p-0101 $ █

```

### 4. Task 6d:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nano task6d.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6d.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
11 22 33 44
00 00 00 00
00 00 00 00
00 00 00 00

pdc-p200045@lmar ~/cuda/20p-0101 $ █

```

### 5. Task 6e:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6e.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
11 00 00 00
22 00 00 00
33 00 00 00
44 00 00 00

```

## 6. Task 6f:

```

Position1 : int index = threadIdx.x * blockDim.x;
array[index] = (index % 5 == 0) ? (index / 5 + 1) * 11 : 0;
Position2: dim3 dimGrid(4,1,1); dim3 dimBlock(4,1,1);

```

## Result:

```

11 00 00 00
00 22 00 00
00 00 33 00
00 00 00 44

```

## 7. Task 6g:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6g.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
11 22 33 44
00 00 00 00
00 00 00 00
00 00 00 00

```

## 8. Task 6f:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6f.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
11 00 00 00
22 00 00 00
33 00 00 00
44 00 00 00

pdc-p200045@lmar ~/cuda/20p-0101 $ 

```

### 9. Task 6g:

```
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6g.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
00 00 00 00
00 00 00 00
00 00 00 00
11 22 33 44

pdc-p200045@lmar ~/cuda/20p-0101 $
```

### 10. Task 6h:

```
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6h.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
00 00 00 11
00 00 00 22
00 00 00 33
00 00 00 44

pdc-p200045@lmar ~/cuda/20p-0101 $
```

### 11. Task 6j:

```
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6j.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
11 22 33 44
11 22 33 44
11 22 33 44
11 22 33 44

pdc-p200045@lmar ~/cuda/20p-0101 $
```

### 12. Task 6k:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nano task6k.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6k.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
11 11 11 11
22 22 22 22
33 33 33 33
44 44 44 44

pdc-p200045@lmar ~/cuda/20p-0101 $ █

```

### 13. Task 6m:

```

Position2: dim3 dimGrid(N/4,1,1) ; dim3 dimBlock(N/4 , 1 , 1);
Position1: int index = threadIdx.x + blockIdx.x * blockDim.x;
array[index] = (4- blockIdx.x) * 11;

```

#### Result:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nano task6m.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6m.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
44 44 44 44
33 33 33 33
22 22 22 22
11 11 11 11

pdc-p200045@lmar ~/cuda/20p-0101 $ █

```

### 14. Task 6n:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nano task6n.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task6n.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
11 11 22 22
11 11 22 22
33 33 44 44
33 33 44 44

pdc-p200045@lmar ~/cuda/20p-0101 $ █

```

### 15. Task 6o:

Position 1:

```
int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < 16) {
        array[i] = (i / 4 + 1) * 11;
    }
```

Position 2:

```
dim3 dimBlock(16); dim3 dimGrid(1);
```

## Task 7: Matrix Addition

```
#include <stdio.h>
#include <stdlib.h>
#define N 16
__global__ void add(int *a, int *b, int *c) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x; // Thread :

    // Position 1: Write Code here for vector addition
    if (tid < N) {
        c[tid] = a[tid] + b[tid];
    }
}

int main() {
    int *a, *b, *c, *da, *db, *dc, i;

    a = (int*)malloc(sizeof(int)*N); // allocate host mem
    b = (int*)malloc(sizeof(int)*N); // and assign random memory
    c = (int*)malloc(sizeof(int)*N); // for the result

    // Write code to initialize both a and b to 1's.
    for (i = 0; i < N; i++) {
        a[i] = 1;
        b[i] = 1;
    }
}
```

```

    cudaMalloc((void **)&da, sizeof(int)*N);
    cudaMalloc((void **)&db, sizeof(int)*N);
    cudaMalloc((void **)&dc, sizeof(int)*N);

    cudaMemcpy(da, a, sizeof(int)*N, cudaMemcpyHostToDevice);
    cudaMemcpy(db, b, sizeof(int)*N, cudaMemcpyHostToDevice);

    dim3 dimGrid(N/8, 1, 1);
    dim3 dimBlock(N/4, 1, 1);

    add<<<dimGrid,dimBlock>>>(da, db, dc);

    cudaMemcpy(c, dc, sizeof(int)*N, cudaMemcpyDeviceToHost);

    for (i = 0; i < N; i++) {
        printf("a[%d] + b[%d] = %d\n", i, i, c[i]);
    }

    // Free allocated memory
    free(a);
    free(b);
    free(c);
    cudaFree(da);
    cudaFree(db);
    cudaFree(dc);

    return 0;
}

```

## Result:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task7d.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
a[0] + b[0] = 2
a[1] + b[1] = 2
a[2] + b[2] = 2
a[3] + b[3] = 2
a[4] + b[4] = 2
a[5] + b[5] = 2
a[6] + b[6] = 2
a[7] + b[7] = 2
a[8] + b[8] = 0
a[9] + b[9] = 0
a[10] + b[10] = 0
a[11] + b[11] = 0
a[12] + b[12] = 0
a[13] + b[13] = 0
a[14] + b[14] = 0
a[15] + b[15] = 0
pdc-p200045@lmar ~/cuda/20p-0101 $ █

```

## Task 8: Matrix Addition Slightly Complicated

```

#include <stdio.h>
#include <stdlib.h>

_global_ void add(int *a, int *b, int *c, int N) {
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    int row = blockIdx.y * blockDim.y + threadIdx.y;

    int index=row*blockDim.x*gridDim.x+col;
    c[index] = a[index] + b[index];
}

int main() {
    int *a, *b, *c, *da, *db, *dc, N = 16, i, j;

    a = (int*)malloc(sizeof(int)*N*N); // allocate host mem
    b = (int*)malloc(sizeof(int)*N*N); // and assign random mem
    c = (int*)malloc(sizeof(int)*N*N); // for the result

```

```

// Write code to initialize both a and b to 1's.
for (i = 0; i < N*N; i++) {
    a[i] = 1;
    b[i] = 1;
}

cudaMalloc((void **)&da, sizeof(int)*N*N);
cudaMalloc((void **)&db, sizeof(int)*N*N);
cudaMalloc((void **)&dc, sizeof(int)*N*N);

cudaMemcpy(da, a, sizeof(int)*N*N, cudaMemcpyHostToDevice);
cudaMemcpy(db, b, sizeof(int)*N*N, cudaMemcpyHostToDevice);

dim3 dimGrid(N/8, N/8, 1);
dim3 dimBlock(N/8, N/8, 1);

add<<<dimGrid, dimBlock>>>(da, db, dc, N);

cudaMemcpy(c, dc, sizeof(int)*N*N, cudaMemcpyDeviceToHost);

for (j = 0; j < N; j++) {
    for (i = 0; i < N; i++) {
        printf("a[%d] + b[%d] = %d\n", j*N+i, j*N+i, c[j*N+i]);
    }
    printf("\n");
}

// Free allocated memory
free(a);
free(b);
free(c);
cudaFree(da);
cudaFree(db);
cudaFree(dc);

```



```
    return 0;
}
```

## Result:

```
pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task8e.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
a[0] + b[0] = 2
a[1] + b[1] = 2
a[2] + b[2] = 2
a[3] + b[3] = 2
a[4] + b[4] = 2
a[5] + b[5] = 2
a[6] + b[6] = 2
a[7] + b[7] = 2
a[8] + b[8] = 2
a[9] + b[9] = 2
a[10] + b[10] = 2
a[11] + b[11] = 2
a[12] + b[12] = 2
a[13] + b[13] = 2
a[14] + b[14] = 2
a[15] + b[15] = 2
```

## Task 9: Measurements

**Title:** To measure anything in CUDA, we can use the following from the CUDA Events API's:

```
#include <stdio.h>
#include <cuda_runtime.h>

//Kernel Function for element-wise addition
__global__ void elementWiseAddition(int *a, int *b, int *result,
    int tid=blockIdx.x*blockDim.x+threadIdx.x;
    if(tid<N){
        result[tid]=a[tid]+b[tid];
    }
}
```

```

}

int main() {
    const int N=16;
    cudaEvent_t start,stop;
    float elapsed;

    //Host arrays
    int h_a[N],h_b[N],h_result[N];
    //Initialize host arrays
    for(int i=0;i<N;i++){
        h_a[i]=i;
        h_b[i]=2*i;
    }
    //Device arrays
    int *d_a,*d_b,*d_result;

    cudaMalloc((void **)&d_a,N*sizeof(int));
    cudaMalloc((void **)&d_b,N*sizeof(int));
    cudaMalloc((void **)&d_result,N*sizeof(int));

    //Copy data from host to device
    cudaMemcpy(d_a,h_a,N*sizeof(int),cudaMemcpyHostToDevice);
    cudaMemcpy(d_b,h_b,N*sizeof(int),cudaMemcpyHostToDevice);
    printf("Grid X\tGrid Y\tGrid Z\tBlock X\tBlock Y\tBlock Z\t\n");
    for(int gridX=1;gridX <= N;gridX*=2){
        for(int gridY=1;gridY<=N/gridX;gridY*=2){
            int gridZ=N/(gridX*gridY);

            for(int blockX=1;blockX<=N;blockX*=2){
                for(int blockY=1;blockY<=N/blockX;blockY*=2){
                    int blockZ=N/(blockX*blockY);

                    dim3 gridSize(gridX,gridY,gridZ);
                    dim3 blockSize(blockX,blockY,blockZ);

```

```

        cudaEventCreate(&start);
        cudaEventCreate(&stop);

        cudaEventRecord(start, 0);

        //call the Kernel
        elementWiseAddition<<<gridSize, blockSize>>>()
        cudaEventRecord(stop, 0);
        cudaEventSynchronize(stop);
        cudaEventElapsedTime(&elapsed, start, stop);

        //Copy result from device to host
        cudaMemcpy(h_result, d_result, N*sizeof(int), cudaMemcpyDeviceToHost);

        // Print the configuration and timing information
        printf("%d\t%d\t%d\t%d\t%d\t%d\t%0.3f\t%0.3f\n",
               gridSizeX, gridSizeY, gridSizeZ, blockSizeX, blockSizeY,
               elapsed * 1000, elapsed, elapsed / 1000);

        // Reset device memory for the next iteration
        cudaMemset(d_result, 0, N * sizeof(int));
    }
}

// Free allocated memory on the device
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_result);

// Destroy CUDA events
cudaEventDestroy(start);
cudaEventDestroy(stop);

```

```

    return 0;
}

```

## Result:

```

pdc-p200045@lmar ~/cuda/20p-0101 $ nvcc task9d.cu
pdc-p200045@lmar ~/cuda/20p-0101 $ ./a.out
Grid X  Grid Y  Grid Z  Block X  Block Y  Block Z  Microseconds  Milliseconds  Seconds
1       1       16       1         1       16       23.616  0.024  0.000
1       1       16       1         2         8       5.760  0.006  0.000
1       1       16       1         4         4       5.856  0.006  0.000
1       1       16       1         8         2       5.856  0.006  0.000
1       1       16       1        16         1       5.728  0.006  0.000
1       1       16       2         1         8       5.728  0.006  0.000
1       1       16       2         2         4       5.760  0.006  0.000
1       1       16       2         4         2       5.760  0.006  0.000
1       1       16       2         8         1       5.760  0.006  0.000
1       1       16       4         1         4       5.760  0.006  0.000
1       1       16       4         2         2       5.728  0.006  0.000
1       1       16       4         4         1       5.760  0.006  0.000
1       1       16       8         1         2       5.760  0.006  0.000
1       1       16       8         2         1       5.952  0.006  0.000
1       1       16      16         1         1       5.792  0.006  0.000
1       2         8         1         1       16       5.824  0.006  0.000
1       2         8         1         2         8       5.792  0.006  0.000
1       2         8         1         4         4       5.824  0.006  0.000
1       2         8         1         8         2       5.728  0.006  0.000
1       2         8         1        16         1       6.208  0.006  0.000
1       2         8         2         1         8       5.760  0.006  0.000
1       2         8         2         2         4       5.728  0.006  0.000
1       2         8         2         4         2       5.760  0.006  0.000
1       2         8         2         8         1       5.792  0.006  0.000
1       2         8         4         1         4       5.760  0.006  0.000
1       2         8         4         2         2       5.760  0.006  0.000
1       2         8         4         4         1       5.760  0.006  0.000
1       2         8         8         1         2       5.760  0.006  0.000
1       2         8         8         2         1       5.760  0.006  0.000
1       2         8        16         1         1       5.760  0.006  0.000
1       4         4         1         1       16       5.760  0.006  0.000
1       4         4         1         2         8       5.792  0.006  0.000
1       4         4         1         4         4       5.760  0.006  0.000
1       4         4         1         8         2       5.728  0.006  0.000
1       4         4         1        16         1       5.760  0.006  0.000
1       4         4         2         1         8       5.728  0.006  0.000
1       4         4         2         2         4       5.824  0.006  0.000
1       4         4         2         4         2       5.728  0.006  0.000
1       4         4         2         8         1       5.856  0.006  0.000
1       4         4         4         1         4       5.760  0.006  0.000
1       4         4         4         2         2       5.792  0.006  0.000
1       4         4         4         4         1       5.760  0.006  0.000
1       4         4         8         1         2       5.760  0.006  0.000
1       4         4         8         2         1       5.760  0.006  0.000
1       4         4        16         1         1       5.888  0.006  0.000
1       8         2         1         1       16       5.760  0.006  0.000
1       8         2         1         2         8       5.920  0.006  0.000

```

---

---

---

---

**END.**